Large feature dot products arise when out-of-sample actions are used in TD-learning compared to SARSA, despite similar Q-values

Large feature dot products <u>eventually</u> correlate with divergent Q-values

